

Enterprise RAG Architecture Guide

1. Chunking: semantic splitting with overlap for better retrieval.
2. Embedding: bge-m3 vectorizes each chunk into 1024 dimensions.
3. Storage: vectors are stored in Qdrant with cosine distance.
4. Retrieval: multi-query recall followed by reranker reranking.
5. Agent: LangGraph orchestrates plan-retrieve-answer pipeline.